

CS310: Advanced Data Structures and Algorithms

Fall 2014 Programming Assignment 6

Due: Thu., December 4, 2014 at midnight.

Goals

This assignment aims to help you:

- Use an open-source Graph package in a graph algorithm
- See how to implement graphs with the help of JDK classes
- Study another example of a type hierarchy at work
- Use Swing support for graph visualization Learn about backtracking and dynamic programming

Description

We will be using the open-source Java graph package “JGraphT”. See its home if you are interested. A subset of its distribution with everything we need is available at `$cs310/jgrapht.zip`. The graph APIs are in package `org.jgrapht`, most significantly `Graph.java`. The `cs310`-related sources are in package `cs310`. You will find javadoc for the whole package, including the `cs310` additions, at `javadoc`. The other packages are exactly as distributed in JGraphT. Some experimental and test-related packages and the big `jgraphmanual.pdf` have been deleted to save on size. Of course you can download the whole thing yourself from the JGraphT site. The `jgraphmanual.pdf` is separately provided in `$cs310`, but it may not be helpful since it covers only the graph display package, not the JGraphT package we will mainly be using. To work on your home computer, download `jgrapht.zip` and install it the same way you did the `games.zip` distribution. For eclipse, set the output directory to build (not classes this time) instead of bin.

JGraphT Graphs

The Graph API in `Graph.java` can be used to describe both directed and undirected graphs, by careful definition of the methods with this flexibility in mind. Concrete classes for directed and undirected graphs are provided, and interfaces `DirectedGraph` and `UndirectedGraph`, with a few methods involving degrees of vertices that can’t be merged together. See `cs310.LoadGraph.java` for an example.

Questions

1. **Topological Sort.** Write a generic `<V, E>` class `TopSort.java` which finds a topological sort for a given directed graph. The constructor takes the `DirectedGraph<V, E>` to work on. Here is the one public method it needs:

```
public List<V> getTopOrder() throws HasCycleException
```

Its output is a List of the vertices of the graph in topological order. Write a client `TopSortTest.java` to test your `TopSort` class without Swing visualization. Like `DFSTest`, `TopSortTest` takes an optional argument for the input file name. With an input file, say `prereq.dat`, in the top-level directory, you would use the command line:

```
cd build
java cs310.TopSortTest ../prereq.dat      (with no arg, uses test.dat)
```

Some sample input files are prereq.dat and cycle.dat described below, as well as test.dat you can find in the JGraphT package (you can ignore the weights of course). Here is the output from those two files:

Input graph

```
110 210
110 240
210 310
240 310
140 310
310 450
450 650
```

topological sort:

```
[110, 210, 240, 140, 310, 450, 650]
```

Input graph

```
B    foo
foo  bar
bar  baz
baz  foo
B    C
X    C
bar  OK
```

topological sort:

```
cycle: [foo, bar, baz, foo]
```

Note that topological sort order is not unique. The second example covers the case that `topSort.getTopOrder()` throws a `HasCycleException` when the input directed graph `g` contains a cycle. Your program should check if a graph has a cycle (and throw an exception if it does) but you DON'T have to determine the cycle itself. Do it only if you want to. If you do determine a cycle, provide another method to return it to the client.

Also write a client `TopSortDemo.java` to display the completed sort using colors as done in `DFS Demo`, that is blue first, etc. `TopSortDemo` needs no arguments on the command line and uses the random graph generator just as `DFS Demo` does.

Here are two ways to implement Topological Sort:

1. Use the method given in Weiss on pages 499–501.
2. Use Depth First Search postorder. See class notes. This should require just a few lines of code.

You can choose any method you want, except using the canned JGraphT algorithms and iterators – that's too easy.

2. **Implementing the Graph API.** We have provided a simple undirected graph implementation in `$cs310/jgrapht/src/cs310/SimpleGraph1.java`, following the general methodology of the JGraphT implementations, but without all the optimizations. The `removeVertex` and `removeEdge` methods are not functional yet, so you have to implement them. Modify `SimpleGraphTest.java` so that it thoroughly exercises your code.
3. **Dijkstra's Algorithm.** Write a class `Dijkstra.java` that finds (as efficiently as possible) the least-cost paths from a given node to all nodes in the given directed graph, given non-negative edge costs as weights in the given `Graph`. Note `LoadWeightedGraph` in the `cs310` directory.

```
// a Dijkstra HASA DirectedGraph
public Dijkstra( DirectedGraph<V,E> g)
// compute best costs from start
public Map<V,Double> dijkstra(V start);
// return least-cost path from start to target
public List<V> pathTo(V start, V target);
```

Remember the start vertex used in a call to `dijkstra` so that you can reuse data structures if `pathTo` is called with the same start vertex. Provide `DijkstraTest.java` so that the command (after `cd build`, and assuming `yourgraph.dat` is in the top-level directory)

```
java cs310.DijkstraTest ../yourgraph.dat
```

will output a table showing the best distances from the first node in the data to all others, and the shortest path from the first node.

A sample input file `test.dat` should produce output like this, in any order:

Shortest Distances and their Paths from A

```
A 0 (A)
B 1 (A, B)
X 3 (A, B, X)
Y 5 (A, B, X, Y)
Z 6 (A, B, X, Y, Z)
C 2 (A, C)
J 5 (A, C, J)
M INFINITY ()
```

Also provide a program `DijkstraDemo.java`, that takes a randomly-generated graph (`UnweightedDemo` uses `LoadGraph`, do a simple modification of that to use `LoadWeightedGraph`) and shows the progress of `dijkstra` vertex by vertex, coloring the vertices of newly-determined distance in the same color sequence defined in `UnweightedDemo.java`. Its command line execution, after `cd` to `build`:

```
java -cp ../lib/jgraph.jar;. cs310.DijkstraDemo
```

For UNIX execution, replace the semicolon above with a colon.

We know of one easy way to implement Dijkstra efficiently, using a trick that allows us to use an ordinary priority queue and yet change the effective priority of elements when we need to. This trick is discussed in the class notes and is coded in Weiss on page 495, but for pa6 be sure to use `JGraphT`.

For the priority queue, set up a `PriorityQueue` and `Path` as in Weiss, but you also need a `Map` (something like the `Map` from `V` to `Unweighted.DistInfo` in `cs310.Unweighted`) to handle the information held in Weiss's "scratch" spot, `v.scratch`, as well as `v.dist` and `v.prev`. But don't use the name "scratch". This is a flag for the vertex having been processed, and deserves a name that says that. The `Path` class in the book is usable as is.

memo.txt

1. Do questions 6 and 7 in HW5. Submit in class on Thursday, Dec. 4.
2. For `SimpleGraph1`, what is the time complexity of methods `addVertex(V)` and `addEdge(V,V)`, in terms of the number `N` of nodes and `E` of edges in the graph?
3. We are using `HashMaps` in our graph implementation. Could we switch over to `TreeMaps` by just replacing "HashMap" with "TreeMap", or would there be problems in doing this?
4. If we were using lots of little graphs, the memory overhead of all these `HashMaps` could be a problem. What is the default initial size of the hash table? How can you halve it?
5. What is the big-Oh performance of your Dijkstra implementation? Show its derivation and what assumptions you are using.
6. Estimate the big-Oh memory usage with `N` nodes and $O(1)$ edges leaving each node.
7. If you used any late days, indicate it in your `memo.txt`.